

# FIT5149: Applied Data Analysis

## Support Vector Machines


Dr. Du, Lan


Faculty of Information Technology, Monash University, Australia


Week 9


- 1 Introduction
- 2 Maximal Margin Classifier
- 3 Support Vector Classifiers
- 4 Support Vector Machines


# Introduction













# Maximal Margin Classifier


## What is a Hyperplane?

- In a  $p$ -dimensional space, a hyperplane is a flat affine subspace of dimension  $p - 1$ 
  - ▶ In 2 dimensional space: a flat one-dimensional subspace is a line.
  - ▶ In 3 dimensional space: a flat two-dimensional subspace is a plane.
  - ▶ More than 3 dimensional space: No visualisation
- $\mathbf{X} = (X_1, X_2, \dots, X_p) \in \mathbb{R}^p$  the  $p$ -dimensional hyperplane is

$$\beta_0 + \beta_1 X_1 + \beta_2 X_2 + \dots + \beta_p X_p = 0$$


- ▶ The vector  $\boldsymbol{\beta} = \beta_1, \beta_2, \dots, \beta_p$  is called the normal vector.
  - It is orthogonal to the surface of a hyperplane.

# What is a Hyperplane?



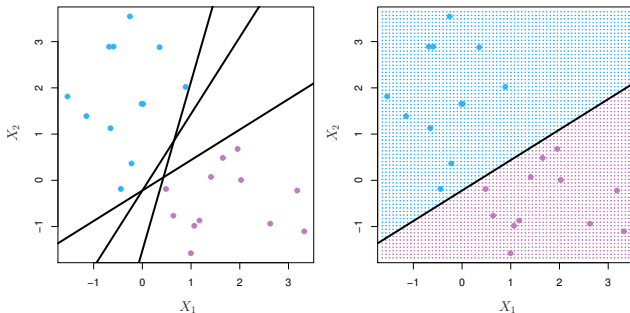
- The magnitude of  $f(X) = \beta_1 X_1 + \beta_2 X_2 - 6$ : the certainty about the class assignment of  $X$ .

# Separating Hyperplane



- If  $f(\mathbf{X}) = \beta_0 + \boldsymbol{\beta}^T \mathbf{X}$ , then  $f(\mathbf{X}) > 0$  for points on one side of the hyperplane, and  $f(\mathbf{X}) < 0$  for points on the other.
- If we code the coloured points as  $Y_i = +1$  for blue and  $Y_i = -1$  for purple, then
  - ▶  $Y_i \times f(\mathbf{X}_i) > 0$  for all  $i$ ,
  - ▶  $f(\mathbf{X}) = 0$  defines a separating hyperplane.
- The magnitude of  $f(\mathbf{X})$  can measure the confidence of the class assignment of  $\mathbf{X}$ .


# Separating Hyperplane



- Decide which of the infinite possible separating hyperplanes to use.


# Maximal Margin Classifier

The idea: find the hyperplane that makes the biggest gap or margin between the two classes.



- 1 Compute the (perpendicular) distance from each training observation to a given separating hyperplane;
- 2 find the minimal distance from the observations to the hyperplane, known as **margin**.
- 3 **Target:** Maximise the margin.


## Functional margin vs geometric margin



Functional margin:  $\hat{M}_i = y_i(\boldsymbol{\beta}^T \mathbf{x}_i + \beta_0)$

- if  $y_i = +1$ , for the functional margin to be large (i.e., for our prediction to be confident and correct), then we need  $\boldsymbol{\beta}^T \mathbf{x}_i + \beta_0$  to be a large positive number.
- if  $y_i = -1$ , for the functional margin to be large, then we need  $\boldsymbol{\beta}^T \mathbf{x}_i + \beta_0$  to be a large negative number.

## Functional margin vs geometric margin



Functional margin:  $\hat{M}_i = y_i(\boldsymbol{\beta}^T \mathbf{x}_i + \beta_0)$

- For a linear classifier with a sign function  $g$ , if we replace  $\boldsymbol{\beta}$  with  $2\boldsymbol{\beta}$  and  $\beta_0$  with  $2\beta_0$ , then

$$g(\boldsymbol{\beta}^T \mathbf{x}_i + \beta_0) = g(2\boldsymbol{\beta}^T \mathbf{x}_i + 2\beta_0)$$


Conclusion: by exploiting our freedom to scale on the parameters, we can make the functional margin arbitrarily large without really changing anything meaningful.

- Given a set of training set  $S = (\mathbf{x}_i, y_i); i = 1, \dots, N$ , the functional margin

$$\hat{M} = \min_{i=1, \dots, N} \hat{M}_i$$



# Functional margin vs geometric margin




Geometric margin:  $M_i = y_i \frac{\beta^T x_i + \beta_0}{\|\beta\|}$

- if  $\|\beta\| = 1$ , the functional margin equals the geometric margin.
- Given a set of training set  $S = (x_i, y_i); i = 1, \dots, N$ , the geometric margin

$$M = \min_{i=1, \dots, N} M_i$$

# Construct Maximal Margin Classifier



Constrained optimisation problem:

$$\text{maximize } M \quad (1)$$

$$\beta_0, \beta_1, \dots, \beta_p$$


$$\text{subject to } \|\boldsymbol{\beta}\|^2 = 1 \quad (2)$$

$$y_i(\beta_0 + \boldsymbol{\beta}^T \mathbf{x}_i) \geq M$$

$$\text{for all } i = 1, \dots, N \quad (3)$$

Now, we can use the Lagrange duality to solve the constrained optimisation problem.

# Construct Maximal Margin Classifier



Constrained optimisation problem:


$$\min_{\beta_0, \beta_1, \dots, \beta_p} \frac{1}{2} \|\boldsymbol{\beta}\|^2 \quad (1)$$

$$\text{s.t. } y_i(\beta_0 + \boldsymbol{\beta}^T \mathbf{x}_i) \geq 1, \quad (2)$$

for all  $i = 1, \dots, N$

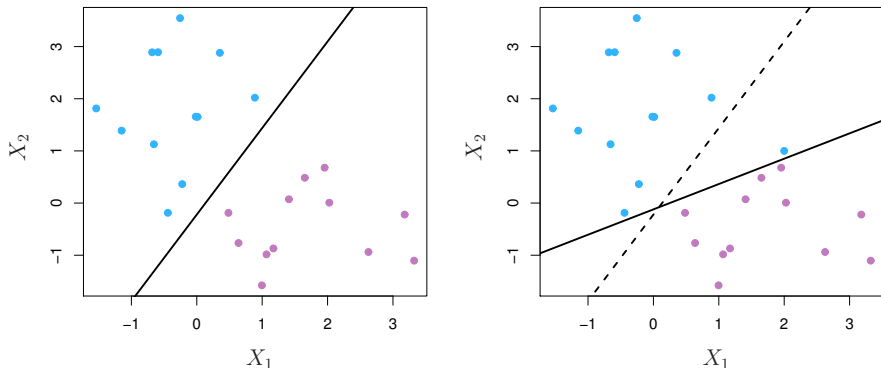
Now, we can use the Lagrange duality to solve the constrained optimisation problem.

# The Non-separable Case



- No separating hyperplane exists
- This often the case, unless  $N < p$

## Sensitivity to Noisy Data





**Figure:** Right: An additional blue observation has been added, leading to a dramatic shift in the maximal margin hyperplane shown as a solid line. The dashed line indicates the maximal margin hyperplane that was obtained in the absence of this additional point.

# Support Vector Classifiers

# Support vector classifier

The idea: use a soft margin



- Soft margin classifier: allow some observations to be on the incorrect side of the margin, or even the incorrect side of the hyperplane

## Construct the soft margin classifier

$$\max_{\beta_0, \boldsymbol{\beta}, \epsilon} M \quad (3)$$

Subject to

$$\|\boldsymbol{\beta}\| = 1 \quad (4)$$

$$y_i(\beta_0 + \beta_1 x_{i1} + \beta_2 x_{i2} + \dots + \beta_p x_{ip}) \geq M(1 - \epsilon_i) \quad \text{for all } i \quad (5)$$

$$\epsilon_i \geq 0, \sum_{i=1}^n \epsilon_i \leq C \quad (6)$$

where

- $\epsilon_i$ : tells us where the  $i$ th observation is located, relative to the hyperplane and relative to the margin.
  - ▶  $\epsilon_i = 0$ :  $\mathbf{x}_i$  is on the correct side of the margin
  - ▶  $\epsilon_i > 0$ :  $\mathbf{x}_i$  is on the wrong side of the margin
  - ▶  $\epsilon_i > 1$ :  $\mathbf{x}_i$  is on the wrong side of the hyperplane



## Construct the soft margin classifier

$$\max_{\beta_0, \boldsymbol{\beta}, \epsilon} M \quad (3)$$

Subject to

$$\|\boldsymbol{\beta}\| = 1 \quad (4)$$

$$y_i(\beta_0 + \beta_1 x_{i1} + \beta_2 x_{i2} + \dots + \beta_p x_{ip}) \geq M(1 - \epsilon_i) \quad \text{for all } i \quad (5)$$

$$\epsilon_i \geq 0, \sum_{i=1}^n \epsilon_i \leq C \quad (6)$$

where

- $C$ : the tuning parameter
  - ▶  $C = 0$ : no budget for violations to the margin.
  - ▶  $C > 0$ : no more than  $C$  observations can be on the wrong side of the hyperplane.

## Construct the soft margin classifier

$$\max_{\beta_0, \boldsymbol{\beta}, \epsilon} M \quad (3)$$

Subject to

$$\|\boldsymbol{\beta}\| = 1 \quad (4)$$

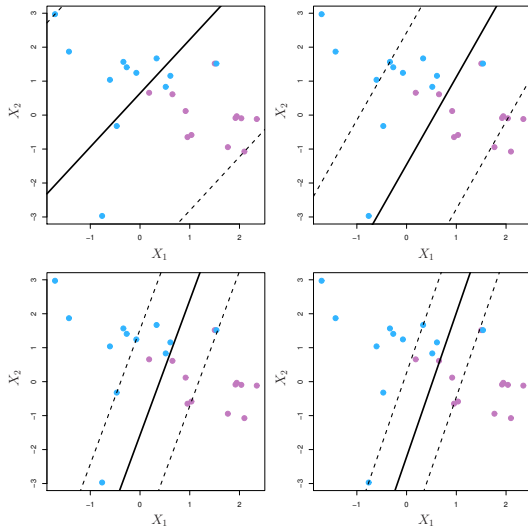
$$y_i(\beta_0 + \beta_1 x_{i1} + \beta_2 x_{i2} + \dots + \beta_p x_{ip}) \geq M(1 - \epsilon_i) \quad \text{for all } i \quad (5)$$

$$\epsilon_i \geq 0, \sum_{i=1}^n \epsilon_i \leq C \quad (6)$$

where



- $C$ : the tuning parameter
  - ▶ the bias-variance trade-off
    - Large  $C$ : wide margin, many support vectors, highly fit to the data, high bias, low variance
    - Small  $C$ : narrow margin, few support vectors, fit the data less hard, low bias, high variance

## C is a regularisation parameter



- Largest value of  $C$  was used in the top left panel, and smaller values were used in the top right, bottom left, and bottom right panels.
- The support vector classifier's decision rule is based only on the support vectors

## Linear boundary can fail



What to do?

# Support Vector Machines

## Feature expansion

- Enlarge the space of features by including transformations;

$$X_1, X_1^2, X_2, X_2^2, \dots, X_p, X_p^2$$

- Fit a support-vector classifier in the enlarged space results in non-linear decision boundaries in the original space.
- Example: Suppose we use  $(X_1, X_2, X_1^2, X_2^2, X_1X_2)$  instead of just  $(X_1, X_2)$ . Then the decision boundary would be of the following form

$$\beta_0 + \beta_1 X_1 + \beta_2 X_2 + \beta_3 X_1^2 + \beta_4 X_2^2 + \beta_5 X_1 X_2 = 0$$

## Feature expansion

- Then the problem is

$$\max_{\beta_0, \beta, \epsilon} M \quad (7)$$


Subject to

$$y_i \left( \beta_0 + \sum_{j=1}^p \beta_{j1} x_{ij} + \sum_{j=1}^p \beta_{j2} x_{ij}^2 \right) \geq M(1 - \epsilon_i) \quad \text{for all } i \quad (8)$$

$$\epsilon_i \geq 0, \quad \sum_{i=1}^n \epsilon_i \leq C, \quad \sum_{j=1}^p \sum_{k=1}^2 \beta_{jk}^2 = 1 \quad (9)$$

$$(10)$$

## Cubic polynomial



- A basis expansion of cubic polynomials
- From 2 variables to 9 variables
- The support-vector classifier in the enlarged space solves the problem in the lower-dimensional space

$$\beta_0 + \beta_1 X_1 + \beta_2 X_2 + \beta_3 X_1^2 + \beta_4 X_2^2 + \beta_5 X_1 X_2 + \beta_6 X_1^3 + \beta_7 X_2^3 + \beta_8 X_1 X_2^2 + \beta_9 X_1^2 X_2$$





# Support Vector Machine (SVM)

- The use of kernels: a more elegant and controlled way to introduce nonlinearities in support-vector classifiers.
  - ▶ An efficient computational approach.
- So,
  - ▶ What are the Kernels?
  - ▶ How do they work?

## Learning SVM: Minimisation

The optimisation problem for finding the optimal margin classifier

$$\min_{\boldsymbol{\beta}, \beta_0} \frac{1}{2} \|\boldsymbol{\beta}\|^2 \quad (11)$$

$$\text{Subject to } y_i(\boldsymbol{\beta}^T \mathbf{x}_i + \beta_0) \geq 1, \quad i = 1, \dots, N \quad (12)$$

Reformulate the constraints as

$$g_i(\boldsymbol{\beta}) = -y_i(\boldsymbol{\beta}^T \mathbf{x}_i + \beta_0) + 1 \leq 0$$

Construct the Lagrangian

$$\mathcal{L}(\boldsymbol{\beta}, \beta_0, \boldsymbol{\alpha}) = \frac{1}{2} \|\boldsymbol{\beta}\|^2 - \sum_{i=1}^N \alpha_i [y_i(\boldsymbol{\beta}^T \mathbf{x}_i + \beta_0) - 1]$$

## Learning SVM: Minimisation

$$\mathcal{L}(\boldsymbol{\beta}, \beta_0, \boldsymbol{\alpha}) = \frac{1}{2} \|\boldsymbol{\beta}\|^2 - \sum_{i=1}^N \alpha_i [y_i(\boldsymbol{\beta}^T \mathbf{x}_i + \beta_0) - 1]$$

- Set the derivative of  $\mathcal{L}$  with respect to  $\boldsymbol{\beta}$  to zero, we have

$$\boldsymbol{\beta} = \sum_i^N \alpha_i y_i \mathbf{x}_i \quad (13)$$

- Set the derivative of  $\mathcal{L}$  with respect to  $\beta_0$  to zero, we have

$$\sum_{i=1}^N \alpha_i y_i = 0 \quad (14)$$

# Learning SVM: Maximization

$$\mathcal{L}(\boldsymbol{\beta}, \beta_0, \boldsymbol{\alpha}) = \sum_{i=1}^N \alpha_i - \frac{1}{2} \sum_{i,j=1}^N y_i y_j \alpha_i \alpha_j \mathbf{x}_i^T \mathbf{x}_j$$

- The dual optimisation problem:

$$\max_{\boldsymbol{\alpha}} \mathcal{L}(\boldsymbol{\beta}, \beta_0, \boldsymbol{\alpha}) \quad (15)$$

subject to

$$\alpha_i \geq 0, \quad i = 1, \dots, N \quad (16)$$

$$\sum_{i=1}^N \alpha_i y_i = 0 \quad (17)$$

## SVM: Inner product

- The inner product from

$$\boldsymbol{\beta}^T \mathbf{x} + \beta_0 = \left( \sum_{i=1}^N \alpha_i y_i \mathbf{x}_i \right)^T \mathbf{x} + \beta_0 \quad (18)$$

$$= \sum_{i=1}^N \alpha_i y_i \langle \mathbf{x}_i, \mathbf{x} \rangle + \beta_0 \quad (19)$$

- To estimate the parameter  $\boldsymbol{\alpha}$  and  $\beta_0$ , all we need are the  $\binom{n}{2}$  inner products.
- It turns out that most of the  $\alpha_i$  can be zero:

$$\sum_{i \in \mathcal{S}} \hat{\alpha}_i y_i \langle \mathbf{x}_i, \mathbf{x} \rangle + \beta_0$$

where  $\mathcal{S}$  is the support set of indices  $i$  such that  $\hat{\alpha}_i > 0$ .

## SVM: Kernels

$$\langle \mathbf{x}_i, \mathbf{x}_{i'} \rangle = \sum_{j=1}^p x_{ij} x_{i'j}$$

can be generalised to

$$K(\mathbf{x}_i, \mathbf{x}_{i'}) = \phi(\mathbf{x}_i)^T \phi(\mathbf{x}_{i'})$$

where, for instance,

$$\phi(x) = \begin{bmatrix} x \\ x^2 \\ x^3 \end{bmatrix}$$

## Kernels: implicit feature mapping

Suppose  $\mathbf{x} \in \mathbb{R}^p$

If  $p = 3$

$$K(\mathbf{x}_i, \mathbf{x}_{i'}) = (\mathbf{x}_i^T \mathbf{x}_{i'})^2 \quad (20)$$

$$= \left( \sum_{j=1}^p x_{ij} x_{i'j} \right) \left( \sum_{j'=1}^p x_{ij'} x_{i'j'} \right) \quad (21)$$

$$= \sum_{j=1}^p \sum_{j'=1}^p x_{ij} x_{ij'} x_{i'j} x_{i'j'} \quad (22)$$

$$= \sum_{j,j'=1}^p (x_{ij} x_{ij'}) (x_{i'j} x_{i'j'}) \quad (23)$$

$$\phi(\mathbf{x}) = \begin{bmatrix} x_1 x_1 \\ x_1 x_2 \\ x_1 x_3 \\ x_2 x_1 \\ x_2 x_2 \\ x_2 x_3 \\ x_3 x_1 \\ x_3 x_2 \\ x_3 x_3 \end{bmatrix}$$

- Calculating the high-dimensional  $\phi(\mathbf{x})$  requires  $\mathcal{O}(p^2)$  time,
- Finding  $K(\mathbf{x}_i, \mathbf{x}_{i'})$  takes only  $\mathcal{O}(p)$  time

## Kernels: Polynomial kernel

$$K(\mathbf{x}_i, \mathbf{x}_{i'}) = (\mathbf{x}_i^T \mathbf{x}_{i'} + c)^d \quad (24)$$

computes the inner-product needed for  $d$  dimensional polynomials.

- In general, the kernel above corresponds to a feature mapping to

$$\binom{p+d}{d}$$

feature space, corresponding to all monomials of the form  $x_{i1}x_{i2} \dots x_{ip}$  that are up to order  $d$ .



## Kernels: Polynomial kernel

If  $p = 3$ ,

If  $d = 2$ ,


$$K(\mathbf{x}_i, \mathbf{x}_{i'}) = (\mathbf{x}_i^T \mathbf{x}_{i'} + c)^2 \quad (25)$$

$$\begin{aligned}
 &= \sum_{j,j'=1}^p (x_{ij}x_{ij'})(x_{i'j}x_{i'j'}) \\
 &\quad + \sum_{j=1}^p (\sqrt{2c}x_{ij})(\sqrt{2c}x_{i'j}) \\
 &\quad + c^2 \quad (26)
 \end{aligned}$$

$$\phi(x) = \begin{bmatrix} x_1x_1 \\ x_1x_2 \\ x_1x_3 \\ x_2x_1 \\ x_2x_2 \\ x_2x_3 \\ x_3x_1 \\ x_3x_2 \\ x_3x_3 \\ \sqrt{2c}x_1 \\ \sqrt{2c}x_2 \\ \sqrt{2c}x_3 \\ c \end{bmatrix}$$



# SVM: Radial Kernel

$$K(\mathbf{x}, \mathbf{x}_{i'}) = \exp\left(-\gamma \sum_{j=1}^p (x_{ij} - x_{i'j})^2\right)$$





- Implicit feature space: very high dimensional.
- Controls variance by squashing down most dimensions severely.

## SVM: Heart Data



ROC curve is obtained by changing the threshold 0 to threshold  $t$  in  $f(X) > t$ , and recording false positive and true positive rates as  $t$  varies. Here we see ROC curves on training data.

# SVM: Heart Data





## More than 2 classes:

- One-Versus-All Classification
  - ▶ Fit  $K$  different 2-class SVM classifiers  $\hat{f}_k(x)$ ,  $k = 1, \dots, K$ ; each class versus the rest. Classify  $x^*$  to the class for which  $\hat{f}_K(x^*)$  is largest.
- One-Versus-One Classification
  - ▶ Fit all  $\binom{K}{2}$  pairwise classifiers  $\hat{f}_{k,l}(x)$ . Classify  $x^*$  to the class that wins the most pairwise competitions.
- Which to choose? If  $K$  is not too large, use OVO.

# Summary

- SVM
  - ▶ "Support Vector Machines", Chapter 9 of "Introduction to Statistical Learning", 6th edition
- Acknowledgement:
  - ▶ Figures in this presentation were taken from "An Introduction to Statistical Learning, with applications in R" (Springer, 2013) with permission from the authors: G. James, D. Witten, T. Hastie and R. Tibshirani
  - ▶ Some of the slides are reproduced based on the slides from T. Hastie and R. Tibshirani
  - ▶ Some of the deduction formulas are adapted from Andrew Ng's note on SVM.